

# CodeScore

Dept. of CIS - Senior Design 2013-2014\*

Allison Pearce  
alpearce@upenn.edu  
Univ. of Penn  
Philadelphia, PA

Spencer Lee  
lesp@upenn.edu  
Univ. of Penn  
Philadelphia, PA

Tanvir Ahmed  
tanvir@upenn.edu  
Univ. of Penn  
Philadelphia, PA

Will McDermid  
wmcd@upenn.edu  
Univ. of Penn  
Philadelphia, PA

## ABSTRACT

*The technology industry lacks automated tools for evaluating software quality. These tools would be helpful for individuals desiring to improve their abilities, recruiters searching for top programmers, and educators needing to quickly assess student performance. This work describes CodeScore, and application that assesses internal software quality by detecting code smells. Code smells are easily recognized design weaknesses that may indicate more significant problems within the system. The program recognizes code smells using a set of specific rules based on the abstract syntax of the source code. It then uses the results to compute a CodeScore: a single value that reflects the maintainability and understandability of the piece of software based on the presence of code smells.*

## 1. INTRODUCTION

Software failures result in annoyed users at best, and they can cause catastrophic system failures at worst. Six people received massive radiation overdoses from a radiation therapy device in one of the canonical examples of a fatal software error [14]. Failures are consequences of poor software quality. Software quality is defined as “conformance to explicitly stated functional and performance requirements, explicitly documented development standards, and implicit characteristics that are expected of all professionally developed software” [20]. In addition to causing failures, disregard for software quality is expensive and inefficient, requiring more dollars and man-hours to maintain, test, and add features to a project than should be necessary. An increasing appreciation of well-designed software has been manifested in the international standard for software quality, ISO/IEC 25010 [12]. The standard affirms the importance of specifying and measuring characteristics of quality, and it defines quality models intended to aid in accomplishing these goals.

Software quality is commonly divided into internal quality and external quality. External quality describes how well the software meets performance expectations, and encompasses characteristics such as correctness, accuracy, integrity, robustness, reliability, usability, adaptability, and efficiency. Internal quality describes how well the software was designed and implemented, and is concerned with maintainability, flexibility, portability, reusability, readability, understandability, and testability [16]. Internal and external quality are closely related, and deficiencies in internal quality lead to deficiencies in external quality. For example, code that

is difficult to understand will also be difficult to adapt to new requirements, and code that cannot be easily tested will likely be incorrect or unreliable.

One way to diagnose internal quality issues is by detecting code smells. Code smells are easily recognized design weaknesses on the surface of the code that may indicate deeper, more significant problems within the system. Individual code smells are directly related to problems with specific aspects of internal quality. Examples of code smells include:

- Long parameter lists. Large numbers of inputs to a method make the code difficult to read, and they make it easy for the programmer to introduce a bug by passing in parameters in the wrong order.
- Deeply nested conditional logic. A sequence of nested conditionals is difficult to understand because the normal path of execution is unclear.
- Message chaining (e.g. `A.getB().getC().getD()`). Message chains force class A to depend unnecessarily on classes B and C in order to get information from class D. This could introduce complications when any of the four classes need to be modified.
- Absence of Javadocs. Lack of documentation detracts from understandability and maintainability.
- Line length. Lines longer than the standard 80 characters make it difficult to read code on smaller monitors.
- Method length. Complex, “do everything” methods are arduous to read and to modify.
- Hard coding. Values that are hard coded in multiple places must be changed individually by the programmer, which often introduces bugs.
- Comment to code ratio. This is another form of documentation that should be present in quality code.

CodeScore tests for the presence of all of the code smells on the preceding list in order to evaluate internal quality. The software focuses on code smells that relate to two elements of internal software quality: maintainability and understandability. These elements are underemphasized in most classroom settings. Students work quickly and haphazardly to meet deadlines and never revisit their programs after turning them in, so there is no need to write maintainable or understandable code. However, understandability and maintainability are critical outside of the classroom,

\*Advisor: Chris Murphy (cdmurphy@cis.upenn.edu).

where teams of developers need to understand code they did not write, and to maintain it for years as the product evolves. CodeScore helps students assess and improve these necessary software engineering skills by providing objective feedback to make them aware of their weaknesses. Computer science teachers can use the application when grading assignments to assess quality without spending hours reading through each student’s source code. It also allows recruiters to identify potential employees who will contribute to their company’s products with a high regard for quality at a much lower cost than most current recruiting practices. Overall, this tool has the potential to help the technology industry train and recruit a strong developer workforce who will write programs that are internally sound.

At the core of the program are specific rules for defining code smells. The rules are used by custom code smell detection algorithms. When a user uploads their code to be analyzed via the web interface, each detector runs on each file in the sample. The results of the detections are fed to a scoring module to compute an internal quality score, the CodeScore, for the piece of software. The CodeScore and detailed reports with suggestions for improvement are returned to the user.

## 2. RELATED WORK

Before describing our methods, we provide background information on the different methods of analyzing code, the standards for measuring code quality, and various services and applications that already strive to provide quantitative analyses of code quality.

### 2.1 Code Analysis Methodologies

The majority of existing applications for measuring code quality rely on well-established code analysis techniques. These techniques involve breaking code down into specific units for processing. Such techniques include parsing the source code into control structures [15], tokens [11], assembly instructions [19], or objects [5]. Once the source code has been parsed into some unit (this process is called static analysis), the attributes of the code, such as reliability, security, efficiency, size, and maintainability can be measured by calculating various metrics using the parsed results. The actual approach to measuring these attributes originated in [3] and later became part of the ISO/IEC 25000 series of standards relating to the quality and evaluation of software [12]. These techniques and standards form the foundation for today’s code quality measurement applications, and they are the basis for our approach as well.

### 2.2 Ohloh

Launched in January of 2006, Ohloh [1] is web service and online community owned by Black Duck Software which provides basic metrics for over 650,000 open source projects containing over 24 billion lines of code located in over 620,000 source code repositories and managed by more than 3,000,000 contributors. These projects are retrieved from source control repositories and analyzed. Metrics such as lines of code, amount of source code comments, and estimated software cost (based on the COCOMO model [4]) are provided, along with commit statistics, project activity, and programming language detection. This data is graphically displayed when one views a project’s information on the site. Ohloh also provides global statistics across all projects for different pro-

gramming languages and contributor statistics for different authors of open source code.

Ohloh primarily focuses on tracking project/contributor activity for large open-source projects. CodeScore focuses more on providing code quality-oriented metrics.

### 2.3 Common Weakness Enumeration

The Common Weakness Enumeration (CWE) [17] is a community-developed list of software weaknesses hosted by the MITRE Corporation. The CWE was developed with the intention of providing:

- a common standard of identifying, mitigating, and preventing software weaknesses.
- a common source of measuring weaknesses for software security tools.
- a common language for describing the various types of software weaknesses that exist in computer architecture, software design, and source code.

CWE supports the discovery of common types of software security flaws such as buffer overflows, handler errors, path-name traversals, and resource management errors (amongst others) within code.

The Preliminary List of Vulnerability Examples for Researchers, or PLOVER [6], was an early form of CWE. PLOVER was the first attempt to take real-world examples of software vulnerabilities and abstract them into common classes of more general vulnerabilities that can arise during the software development process. The goal of PLOVER was to make this information available to developers, researchers, and analysts so that they may use it with the goal of improving code security. CWE expands upon PLOVER by establishing community-developed definitions and descriptions of these common weaknesses.

CWE primarily strives to provide standards relating to the weaknesses of code in terms of security. This may be a valuable resource as we develop additional features for CodeScore, as many security weaknesses are indicative of, or are themselves, code smells, but our project has a wider scope. CWE also incorporates community feedback in developing definitions of weaknesses, which is a strategy we also intend to use in verifying and fine-tuning our models.

### 2.4 HIST

Historical Information for Smell deTecton, or HIST [18], is an approach developed to detect five specific code smells using change history from version control systems. The developers of HIST point out that not all code smells are possible to detect using just source code because some are by definition characterized by how the code changes during the project’s development. One example is parallel inheritance hierarchies, in which “every time you make a subclass of one class, you have to make a subclass of another” [9]. Though revision histories often display changes at a file-level granularity, they use a tool called the Change History Extractor to parse changes at a method- and class-level granularity, and then they identify code smells from the parsed logs using specific rules. CodeScore will soon make use of a similar strategy for incorporating version control information into several metrics, but CodeScore’s analysis includes many more metrics than just those that can be detected using revision logs.

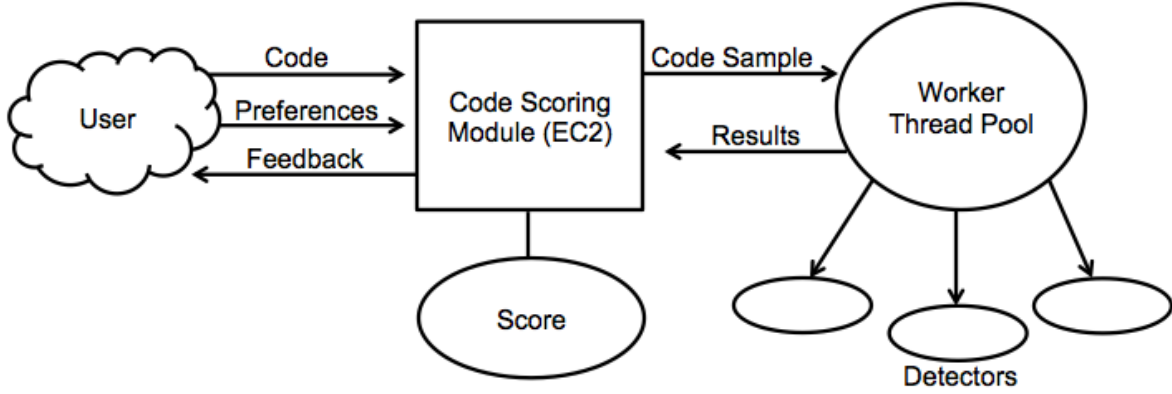


Figure 1: System architecture

## 2.5 JOCS

Judgment of Code Style (JOCS) was a CIS 400/401 project in the 2012-2013 academic year [7]. Their goal was to create a tool for an automated analysis of code style. Such a tool would be useful in assigning the style grade often associated with assignments in introductory programming classes, such as CIS 110 or 120 at the University of Pennsylvania. Features of interest included line length, modularity, and consistency. They used machine learning to compute a single score from the features. CodeScore will make use of similar techniques, but the two projects will be using mostly non-overlapping feature sets and different strategies for identifying them. Our focus is not on style but on code smells as they relate to internal quality, which ultimately affects the correctness, efficiency, and overall external quality of the program.

## 2.6 Other Static Code Analysis Tools

A wide variety of other static code analysis tools also exist, including CodeSonar [10], KlocWork Insight [13], FindBugs [21], and Yasca [22]. These applications focus mostly on external quality metrics such as security, reliability, correct functionality, and efficiency. Most can be used with a variety of programming languages and offer reports or visualizations of the results.

CodeScore offers a novel solution amongst all of these existing tools. While most of the current applications are principally concerned with external quality, our tool focuses on internal quality. Internal quality has a strong correlation with external quality, so writing code with an emphasis on internal quality allows developers to spend less time testing and fixing bugs and more time on new products and features. Additionally, many existing tools target industry developers, but one of our primary target audiences is students. Students need to develop an appreciation for internal quality before entering the workforce or academia, but these skills are often underemphasized at the undergraduate level. Finally, our solution attempts to provide more functionality than the majority of similar open-source applications. CodeScore aims to provide quality of functionality similar to commercial-grade applications in a form that is accessible to a wider audience.

## 3. SYSTEM MODEL

CodeScore focuses on evaluating internal quality, and as such primarily uses static code analysis techniques. This allows for the assessment of source code without bias from different system architectures or environments. Code smells provide a powerful but simple indicator of internal quality because they are specific and detectable, and they have a clear relationship with internal quality. The program currently focuses on understandability and maintainability.

*Understandability* is a measure of how easy a code sample is for a human to interpret. It can be estimated in part by the length of message chains (a code smell in which one method invokes another, which invokes another, and so on in a long one-line sequence of function calls), length of parameter lists, by determining what fraction of variable names are dictionary words versus strings of letters and numbers, and by analyzing the class structure of a program.

*Maintainability* is a measure of how easy a code sample will be to update and change. It can be estimated in part by detecting and recognizing coupling between classes, duplicated code, classes and methods that are too large, and hard coding.

CodeScore implements the workflow illustrated in Figure 1. The key components of the system are the main driver and thread pool, the detectors, and the scoring module. CodeScore includes detectors for the following code smells:

- Long parameter lists. This smell is detected by identifying method declarations and counting the number of parameters in the signature. The maximum allowable parameter list length can be configured by the user.
- Deeply nested conditional logic. This smell is detected by identifying conditional statements and determining if they contain additional conditional statements within the blocks of code that are entered if a particular condition is met. The level of nesting considered acceptable can be configured by the user.
- Message chaining. This is detected by identifying all method invocations and determining if the calling object of the method is itself a method invocation. Note that it is not sufficient simply to detect methods that

appear in the same line, because then an expression such as:

```
System.out.println("x is: " + x.get() + " and y  
is: " + y.get());
```

would be incorrectly labeled a message chain.

- Absence of Javadocs. This is detected by identifying Javadocs and comparing the ratio of the number of methods and classes declared in the software to the number of Javadocs.
- Line length. This is detected by counting the number of characters in each line.
- Method length. This is detected by counting the number of lines of code (not including white space or comments) in a method declaration.
- Hard coding. This is detected by identifying literals (for example, string literals or integer literals). The detector then determines if the literals are being used in an acceptable context, such as a variable assignment, annotation, or in a statement such as `return true;`, or if they are hard-coded values.
- Comment to code ratio. This is detected by differentiating between lines of comments, lines of code, and blank lines, then counting the instances of each.

The detectors are the most complicated component of the system and present the greatest technical challenges. Each code smell requires a unique detection algorithm, sometimes involving analysis of multiple classes together or complex parsing. The detectors search for all possible patterns indicative of a particular code smell using the syntactical structure of the program. They track the number and location of all violations, compiling them into comprehensive reports in javascript object notation (JSON) for interpretation by other modules in the system. In detectors which require thresholds, such as the maximum depth allowed for nested conditionals, the program's default parameters can be overridden by inputting custom parameters.

Users access CodeScore using a web application that provides a platform for developers to assess and showcase their abilities and for recruiters to identify promising candidates. Developers can create an account and upload their projects for analysis. Processing occurs with the detectors running in parallel, and once completed, a report provides the overall CodeScore for the project, which is a grade out of 100 points that is analogous to a grade on an essay. It also shows the subscores of each individual file in the project, graphs displaying the frequency of violations for each metric, and specific suggestions for improvement that include line numbers and snippets from the user's code. This report can be shared as a quick and objective evaluation of one's programming ability, used as a tool for self-improvement, or incorporated into teachers' grading rubrics. The app also charts CodeScores for revisions of projects that are uploaded for re-assessment so that the user can track improvement over time.

The interface for recruiters focuses on searching for and finding information about talented programmers. When viewing a candidate, recruiters can see the candidate's average CodeScore and ranking as compared to other users.

They can also view the candidate's education history, previous employers, answers to sample behavioral questions, and a project showcase.

## 4. SYSTEM IMPLEMENTATION

All processing for CodeScore happens on an Amazon EC2 cloud processing machine. EC2 is a service that provides resizable computing capacity in the cloud. Users upload Java code in the form of a single .java file or a zip file under 10MB in size containing one or more .java files and a preferences file to the EC2 instance through our website, which we developed in PHP. The software ensures that the uploaded project is not malicious using basic validation techniques that search for traits that exist in malware (e.g. known checksums or invalid filetypes). After passing validation, the upload is indexed on Amazon S3. S3 stands for Simple Storage Service and is a component of the Amazon AWS suite of web services. After indexing, the server launches a process that scans for all appropriate files (Java files or preferences files) and saves some of the metadata to include in the report. The preferences file is encoded in JSON and includes information that is specific to each detector, such as thresholds for certain code smells. It also includes optional weights if the user would like to change how much a given detector influences the overall CodeScore. A sample preferences file can be seen in Figure 2. If no preferences file is included, the program only runs using the default values. If a preferences file is included, the program runs once with the default settings to compute a standard CodeScore used to compare the user with other registered candidates and once with the custom settings.

```
{  
  "preferences": {  
    "longestMethodChain": 1,  
    "maxConditionalDepth": 1,  
    "maxParamCount": 3  
  },  
  "weights": {  
    "methodChain": ".21",  
    "nestedConditionals": ".49",  
    "paramCounts": ".3"  
  }  
}
```

Figure 2: User preferences file encoded in JSON

When software is uploaded to the endpoint, the controller class generates  $n$  worker threads, where  $n$  is the number of detection algorithms to be performed on the software sample. Currently,  $n = 8$  because the program supports 8 different detectors. Each of the workers performs a specific detection task on each uploaded file and then sends results back to the scoring module. The type of result reported to the scoring module depends on the detector, but in most cases the detector reports a count of the number of violations of that metric in each file normalized to a percentage between zero and one hundred. The percentage represents how acceptable the file is with regard to that particular code smell, taking into account the length of the file and the number of violations.

The scoring module combines all of the normalized de-

tector outputs into a weighted average representing the subscore for a particular file. Some detectors are weighted more heavily than others. These weights are based on experiments and user surveys. The subscores for each file are combined into the final CodeScore. Each file is weighted by importance, considering factors such as the number of lines it contributed to the project and the number of times methods or objects from that file are used throughout the project. The end result of this process is a single score that gives a user an at-a-glance measure of the internal quality of a project. The web application also provides more detailed feedback using HighCharts [2], a JavaScript framework used for making graphs online. Users can navigate to view reports for a specific file or a general overview.

The program controller, scoring module, and all detectors are implemented in Java. Multithreading and concurrency of the detectors is accomplished using Java’s threading libraries and a custom worker thread pool. The detection algorithms use the Eclipse Java Abstract Syntax Tree (AST) API to parse the Java code as a first step in finding code smells. Traversing the AST allows the system to access and manipulate Java code at the syntactic level, efficiently searching for specific elements such as if statements or method invocations. The AST API also provides information about the location of a particular expression within the tree, both in terms of line numbers and in terms of the surrounding code. Though CodeScore only supports Java at this time, detecting code smells at the AST level makes the program easy to adapt to additional programming languages in the future. Most object-oriented code smells have the same properties from the perspective of the AST, which abstracts away the specifics of individual programming languages. A sample abstract syntax tree and the Java code it represents can be seen in Figure 3.

The web application (techruit.me) is implemented in PHP and makes use of search engine optimization techniques. Custom PHP scripts generate page content, comment, and optimize formatting for each view on the website. The website is also integrated with social media using Facebook’s Open Graph tags. This allows the page to be liked and shared on Facebook and provides information to other users when the page receives attention. Several screenshots of the web application can be seen in Appendix A.

Efficiency is a challenge that is critical to the usefulness of the CodeScore platform. CodeScore is designed to run on large and small software projects, which means that all parsing, processing, and reporting must be carefully engineered to produce results in a timely manner. CodeScore employs parallel processing techniques so that the runtime scales well with the number of metrics and the size of the input.

One drawback of the overall CodeScore system is that it is largely centralized. If the server containing the controller code is compromised at any time during the computation, all progress in the analysis will be lost. A possible future development is to decentralize the process so that loss of a server does not result in complete loss of progress. Ideally, this could be implemented using MapReduce [8].

## 5. RESULTS

CodeScore has been evaluated on four main criteria:

- How accurately code smells are detected.

- How the CodeScore compares to human grading.
- How quickly the analysis is completed.
- The success of the website.

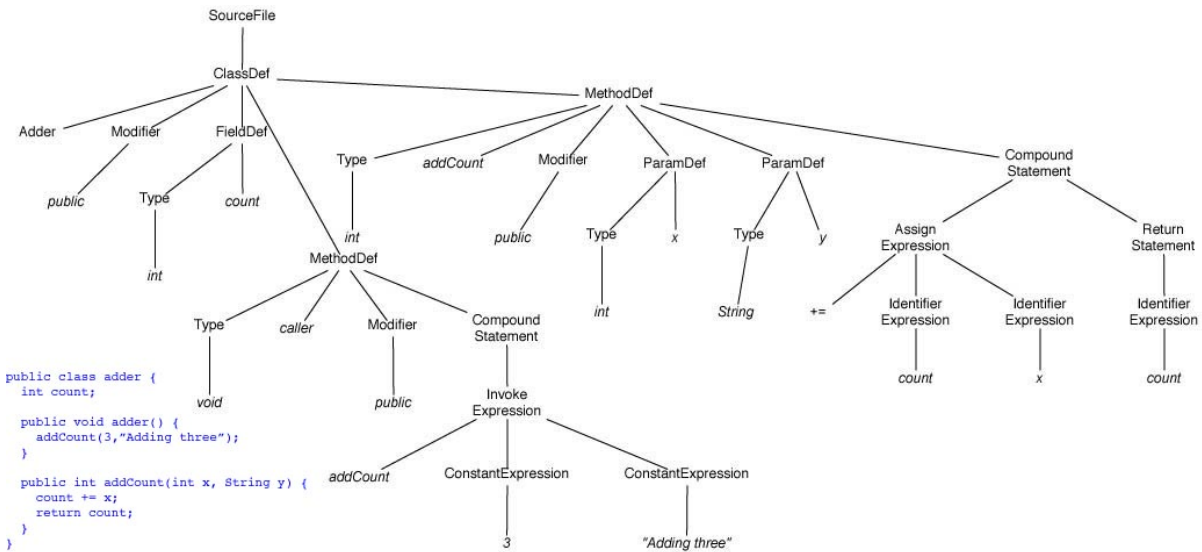
For the first criterion, we designed user surveys in which users identified the same code smells that are detected by CodeScore. For testing small projects, we asked three volunteers to write short Java programs (100-300 lines of code) that included some examples of what they considered “bad code.” For medium-sized projects (500-1000 lines of code), we used assignments submitted by randomly selected students in CIS 120 at the University of Pennsylvania. Finally, to test larger projects, we used open-source Java projects publicly hosted on Github (>2000 lines of code). Each participant evaluated one code sample of each size in three separate sessions for a total of nine code samples per person. Participants were given a list of the code smells, their definitions, and examples, and asked to identify all instances of them in the provided code samples. We considered a code smell “detected” by the participants if at least one third of them (9 people) marked the same line. Figure 4 displays the percentages of code smells correctly identified by the CodeScore detectors as compared to human identification in the user studies.

CodeScore detected 100% of violations using the comment to code ratio, line length, method length, and number of parameters detectors. The absence of Javadocs and message chaining detectors were both over 80%. The detectors for nested conditionals and hard coding are the most technically challenging, which is reflected in their accuracy performance (72% and 68%, respectively). False positives were not an issue with any detectors except for hard coding.

To compare CodeScore’s final grade to the grade a human would assign, we also asked study participants to assign a score from 0 to 100 to each piece of code they analyzed. We used feedback from the first session to fine-tune the weights of the detectors used in calculating the CodeScore. We found that the CodeScore assigned by the software was within 14% of the average score assigned by the study participants for five of the six code samples assessed in the second and third sessions.

To analyze runtime, we pulled additional open source projects with up to 10,000 lines of code to use as inputs to the system. Projects of this size were not feasible to ask user study participants to analyze in the detector and scoring accuracy studies, but they were informative for collecting runtime information. The speed of execution varied roughly linearly with the number of lines of code in preliminary tests. The bottleneck in runtime seems to be the Java Abstract Syntax Tree API. However, the benefits of using the AST API in terms of simplicity and modularity outweigh the slight slowness in processing. See Figure 5 for a graph of the runtime analysis results.

The success of the web application can be quantified using search results and various analytics provided by Google. At the time of publishing of this report, the website (techruit.me) ranks in the top two results when searching Google for “techruit”. The website is in the top 8 results on Google and top 4 on Bing when the search term is “codescore recruiting.” The landing page received over 2,000 pageviews during the time of development, which is highly promising considering that there was no formal marketing. The bounce rate, or the



**Figure 3: Sample Java abstract syntax tree**

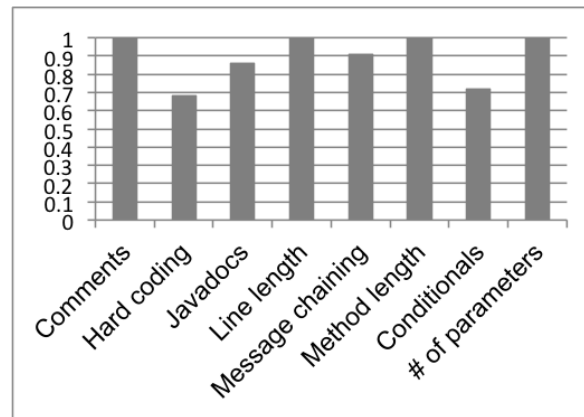


Figure 4: Detector Accuracy


number of users who are interested enough to view other pages after the landing page, is about 50%. About 10.5% of users continue to look at the developers portal and 8.3% look at the recruiters portal.

## 6. FUTURE WORK

Moving forward, there are a number of interesting questions to be investigated with CodeScore. An obvious improvement is to add support for other programming languages in addition to Java. Because CodeScore uses the Java Abstract Syntax Tree API to parse the source code, most of the language-specific considerations have been abstracted away. The detectors look for patterns that are considered detrimental to internal quality in any object-oriented language. Therefore, adding new programming languages that have existing APIs for their syntax trees would simply be a matter of plugging in the new API. A more difficult problem is to include other programming paradigms, such as functional programming.

Even within the realm of Java and object-oriented programming, there are enhancements that could help Code-

Score perform a more thorough assessment of internal quality. Certain code smells, such as shotgun surgery, cannot be detected using just the source code. Shotgun surgery describes the situation in which making one change in behavior requires modifying the code in several places. For example, if logging statements are implemented separately in each function in a class, then adding line numbers to the logs will require considerable time and effort. A better solution would be to write a log wrapper for all of the functions, so that any changes only need to be made once. Using revision history, detecting code smells like shotgun surgery is possible. Analyzing revision history requires a custom parser that can extract changes at method-level granularity in order to determine which methods are changing together. Using revision history might also make it possible to analyze collaborative projects and score each contributor based only on the code that he or she wrote. After creating such a parser and adding the corresponding analyses to the program, it would be simple to link the CodeScore web application to Github and pull entire repositories, including git logs, for processing.



**Figure 5: Preliminary Runtime Analysis of CodeScore Module**

## 7. ETHICS

CodeScore raises few ethical questions. The most significant ethical issue with CodeScore is the possibility of misleading potential employers. If the software were to be adopted by companies to aid in their hiring procedures, and if flaws in the system made it appear that a particular programmer was less qualified than he or she really was, that person might not be hired. CodeScore would be at least partially responsible for keeping that individual from finding work. It is also possible that the opposite could occur - CodeScore might lead to someone being hired despite being underqualified. If the person broke something critical, such as the Apple developer who introduced the “goto fail” bug, CodeScore might also be partially responsible. Other than these scenarios or a similar situations, however, the software presents no major ethical issues. There is almost no physical, financial, or emotional risk to users.

## 8. CONCLUSIONS

CodeScore is a program that performs an automated assessment of the internal quality of a piece of software. The code is graded based on the presence of specific code smells that detract from the maintainability and understandability of software. Code smells are detected using custom algorithms, and the results of these algorithms are used to compute a single numerical score (the “CodeScore”) for the software. The score provides a quick, easy-to-understand measure of the internal quality of the project. The program also provides detailed reports and suggestions for improvement. Of the eight detectors, six achieved over 80% accuracy, and the CodeScore itself was shown to be highly correlated with grades assigned by participants of user studies. CodeScore is useful for individuals who want to improve their abilities, recruiters seeking an inexpensive way to assess potential hires, and instructors who want to reinforce the importance of writing quality code. CodeScore is part of an online platform that allows developers to assess and optionally showcase their code and helps connect recruiters with talented programmers.

## 9. REFERENCES

- [1] J Allen, S Collison, and R Luckey. Ohloh web site. <https://www.ohloh.net/>, 2009.
- [2] HighSoft AS. Highcharts. <http://www.highcharts.com>, 2013.
- [3] Barry W Boehm, John R Brown, and Myron Lipow. Quantitative evaluation of software quality. In *Proceedings of the 2nd international conference on Software engineering*, pages 592–605. IEEE Computer Society Press, 1976.
- [4] Barry W Boehm, Ray Madachy, Bert Steece, et al. *Software Cost Estimation with Cocomo II with Cdrom*. Prentice Hall PTR, 2000.
- [5] Shyam R Chidamber and Chris F Kemerer. A metrics suite for object oriented design. *Software Engineering, IEEE Transactions on*, 20(6):476–493, 1994.
- [6] Steve Christey. Plover: Preliminary list of vulnerability examples for researchers. In *NIST Workshop Defining the State of the Art of Software Security Tools*, 2005.
- [7] Nathan Close, Amalia Hawkins, Sworupini Sureshkumar, et al. Judgement of code style, 2013.
- [8] Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.
- [9] M Fowler. *Refactoring: Improving the design of existing code*. Addison-Wesley, 1999.
- [10] GrammaTech. Codesonar. <http://www.grammatech.com/codesonar>, 2013.
- [11] Maurice H Halstead. *Elements of Software Science*. Elsevier Science Inc., 1977.
- [12] ISO ISO. Iec 25010: 2011: Systems and software engineering—systems and software quality requirements and evaluation (square)—system and software quality models. *International Organization for Standardization*, 2011.
- [13] Klocwork. Klocwork insight. <http://www.klocwork.com/products/insight>, 2013.
- [14] Nancy Leveson. *Safeware: System Safety and Computers*. Addison-Wesley, 1995.
- [15] Thomas J McCabe. A complexity measure. *Software Engineering, IEEE Transactions on*, (4):308–320, 1976.
- [16] Steve McConnell. *Code Complete*. Microsoft Press, 1993.
- [17] CWE MITRE. Common weakness enumeration. <https://cwe.mitre.org/>, 2006.
- [18] Fabio Palomba, Gabriele Bavota, Massimiliano Di Penta, Rocco Oliveto, Andrea De Lucia, and Denys Poshyvanyk. Detecting bad smells in source code using change history information. <http://www.cs.wm.edu/~denys/pubs/ASE%2713-HIST-Camera.pdf>.
- [19] Robert E Park. Software size measurement: A framework for counting source statements. Technical report, DTIC Document, 1992.
- [20] R Pressman. *Software Engineering: A Practitioner’s Approach*. McGraw-Hill, 4 edition, 1997.
- [21] Bill Pugh and Andrey Loskutov. Findbugs. <http://findbugs.sourceforge.net>, 2013.
- [22] Michael Scovetta. Yasca. <http://www.scovetta.com/yasca.html>, 2007.



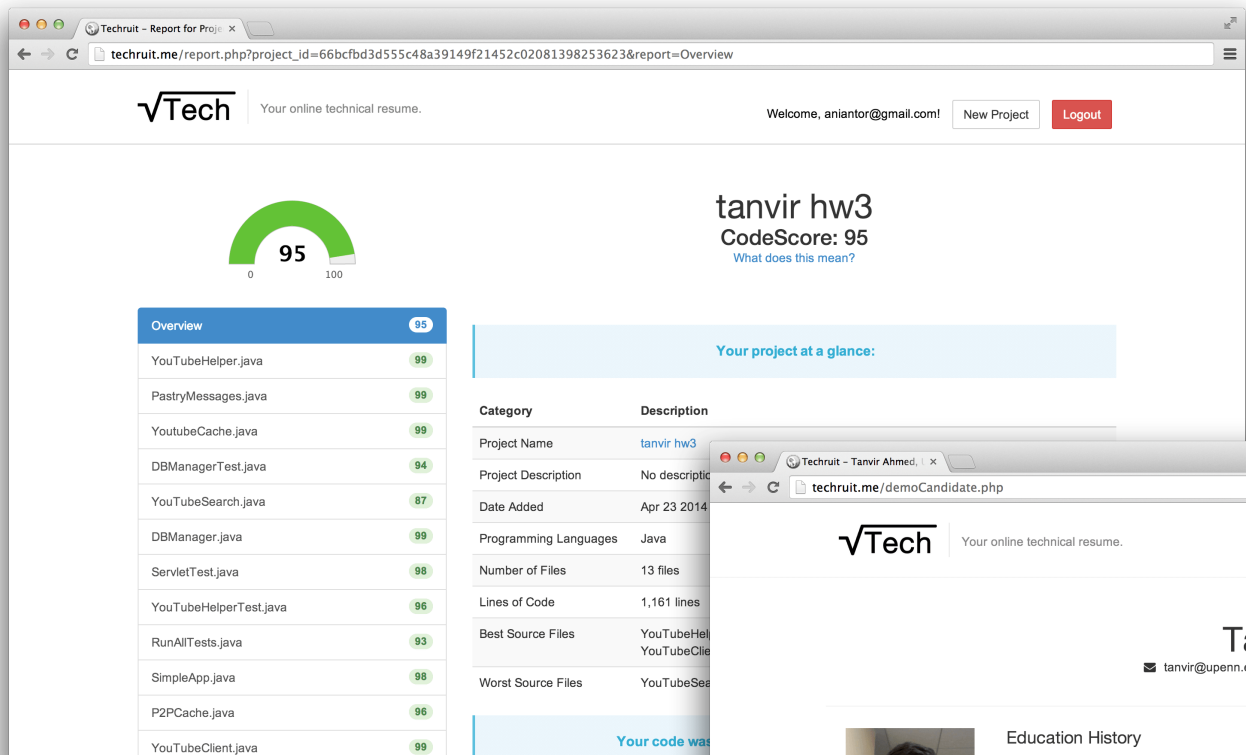


Figure 6: Top of Report Page

## APPENDIX

### A. WEB APPLICATION SCREENSHOTS

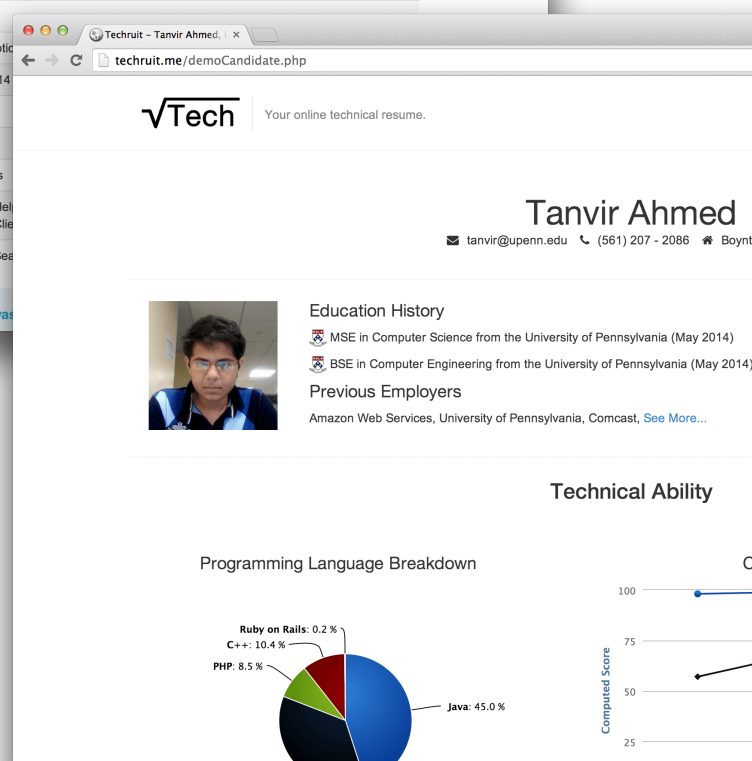


Figure 7: Sample Candidate Profile