CodeScore Progress Report

Dept. of CIS - Senior Design 2013-2014*

Allison Pearce
alpearce@upenn.edu
Univ. of Penn
Philadelphia, PA

Spencer Lee
lesp@upenn.edu
Univ. of Penn
Philadelphia, PA

Tanvir Ahmed
tanvir@upenn.edu
Univ. of Penn
Philadelphia, PA

Will McDermid
wmcd@upenn.edu
Univ. of Penn
Philadelphia, PA

ABSTRACT

The technology industry lacks automated tools for evaluating software quality. These tools would be helpful for individuals desiring to improve their abilities, recruiters searching for top programmers, and educators needing to quickly assess student performance. We are in the process of developing an application called CodeScore that assesses internal software quality by detecting code smells. Code smells are easily recognized design weaknesses that may indicate more significant problems within the system. The program will recognize code smells using a set of specific rules. It will then use the results to compute a CodeScore: a single value that reflects the maintainability and understandability of the piece of software based on the presence of code smells.

1. INTRODUCTION

Software failures result in annoyed users at best, and they can cause catastrophic system failures at worst. Six people received massive radiation overdoses from a radiation therapy device in one of the canonical examples of a fatal software error [13]. Failures are consequences of poor software quality. Software quality is defined as "conformance to explicitly stated functional and performance requirements, explicitly documented development standards, and implicit characteristics that are expected of all professionally developed software" [20]. In addition to causing failures, disregard for software quality is expensive and inefficient, requiring more dollars and man-hours to maintain, test, and add features to a project than should be necessary. An increasing appreciation of well-designed software has been manifested in the international standard for software quality, ISO/IEC 25010. The standard affirms the importance of specifying and measuring characteristics of quality, and it defines quality models intended to aid in accomplishing these goals [11].

Software quality is commonly divided into internal quality and external quality. External quality describes how well the software meets performance expectations, and encompasses characteristics such as correctness, accuracy, integrity, robustness, reliability, usability, adaptability, and efficiency. Internal quality describes how well the software was designed and implemented, and is concerned with maintainability, flexibility, portability, reusability, readability, understandability, and testability [15]. Internal and external quality are closely related, and deficiencies in internal quality lead to deficiencies in external quality. For example, code that is difficult to understand will also be difficult to adapt to

new requirements, and code that cannot be easily tested will likely be incorrect or unreliable.

One way to diagnose internal quality issues is by detecting code smells. Code smells are easily recognized design weaknesses on the surface of the code that may indicate deeper, more significant problems within the system. Some common code smells include:

- Classes or methods that have too many lines of code

- Duplicated code

- Long lists of parameters for functions

- Overly complex conditional logic

- Inconsistent or uncommunicative naming conventions

- Feature envy, a problem in which methods in one class repeatedly use data and methods from another class

Individual code smells are directly related to problems with specific aspects of internal quality. For example, overly complex conditional logic and inconsistent naming conventions detract from understandability and readability. Duplicated code implies that changes will have to be made not once but everywhere that the duplication occurs, which detracts from maintainability.

CodeScore tests for the presence of code smells in order to evaluate internal quality. We have defined rules for custom algorithms to detect code smells, and the results of the detection are used compute an internal quality score for a code sample. The program focuses specifically on how code smells relate to two elements of internal software quality, maintainability and understandability. These elements are underemphasized in most classroom settings. Students work quickly and haphazardly to meet deadlines and never revisit their program after turning it in, so there is no need to write maintainable or understandable code. However, understandability and maintainability are critical outside of the classroom, where teams of developers need to understand code they did not write, and to maintain it for years as the product evolves. CodeScore will help students assess and improve these necessary software engineering skills by providing objective feedback to make them aware of their weaknesses. Computer science teachers can use the the application when grading assignments to assess quality without spending hours reading through each student's source code. It will also allow recruiters to identify potential employees who will contribute to their company's products with a high regard for quality. Overall, this tool will help the technology

---

*Advisor: Chris Murphy (cdmurphy@cis.upenn.edu).

industry to train and recruit a strong developer workforce who will write programs that are internally sound.

2. RELATED WORK

Before describing our methods, we will provide background information on the different methods of analyzing code, the standards for measuring code quality, and various services and applications that already strive to provide quantitative analyses of code quality.

2.1 Code Analysis Methodologies

The majority of existing applications for measuring code quality rely on well-established code analysis techniques. These techniques involve breaking code down into specific units for processing. Such techniques include parsing the source code into control structures [14], tokens [10], assembly instructions [19], or objects [4]. Once the source code has been parsed into some unit (this process is called static analysis), the attributes of the code, such as reliability, security, efficiency, size, and maintainability can be measured by calculating various metrics using the parsed results. The actual approach to measuring these attributes originated in [2] and later became part of the ISO/IEC 25000 series of standards relating to the quality and evaluation of software [11]. These techniques and standards comprise the foundation for today's code quality measurement applications, and they are the basis for our approach, as well.

2.2 Ohloh

Launched in January of 2006, Ohloh [1] is web service and online community owned by Black Duck Software which provides basic metrics for over 650,000 open source projects containing over 24,000,000,000 lines of code located in over 620,000 source code repositories and managed by more than 3,000,000 contributors. These projects are retrieved from source control repositories and analyzed. Metrics such as lines of code, amount of source code comments, and estimated software cost (based on the COCOMO model [3]) are provided, along with commit statistics, project activity, and programming language detection. This data is graphically displayed when one views a project's information on the site. Ohloh also provides global statistics across all projects for different programming languages and contributor statistics for different authors of open source code.

Ohloh primarily focuses on tracking project/contributor activity for large open-source projects. CodeScore focuses more on providing code quality-oriented metrics.

2.3 Common Weakness Enumeration

The Common Weakness Enumeration (CWE) [16] is a community-developed list of software weaknesses hosted by the MITRE Corporation. The CWE was developed with the intention of providing:

- a common standard of identifying, mitigating, and preventing software weaknesses.

- a common source of measuring weaknesses for software security tools.

- a common language for describing the various types of software weaknesses that exist in computer architecture, software design, and source code.

CWE supports the discovery of common types of software security flaws such as buffer overflows, handler errors, pathname traversals, and resource management errors (amongst others) within code.

The Preliminary List of Vulnerability Examples for Researchers (PLOVER) [5] was an early form of CWE. PLOVER was the first attempt to take real-world examples of software vulnerabilities and abstract them into common classes of more general vulnerabilities that can arise during the software development process. The goal of PLOVER was to make this information available to developers, researchers, and analysts so that they may use it with the goal of improving code security. CWE expands upon PLOVER by establishing community-developed definitions and descriptions of these common weaknesses.

CWE primarily strives to provide standards relating to the weaknesses of code in terms of security. This may be a valuable resource as we develop additional features for CodeScore, as many security weaknesses are indicative of, or are themselves, code smells, but our project has a wider scope. CWE also incorporates community feedback in developing definitions of weaknesses, which is a strategy we also intend to use in verifying and fine-tuning our models.

2.4 HIST

Historical Information for Smell deTection, or HIST [18], is an approach developed to detect five specific code smells using change history from version control systems. The developers of HIST point out that not all code smells are possible to detect using just source code because some are by definition characterized by how the code changes during the project's development. One example is parallel inheritance hierarchies, in which "every tme you make a subclass of one class, you have to make a subclass of another" [8]. Though revision histories often display changes at a file-level granularity, they use a tool called the Change History Extractor to parse changes at a method- and class-level granularity, and then they identify code smells from the parsed logs using specific rules. CodeScore will soon make use of a similar strategy for incorporating version control information into several metrics, but CodeScore's analysis includes many more metrics than just those that can be detected using revision logs.

2.5 JOCS

Judgment of Code Style (JOCS) was a CIS 400/401 project in the 2012-2013 academic year [CITE THIS SHIT]. Their goal was to create a tool for an automated analysis of code style. Such a tool would be useful is assigning the style grade often associated with assignments in introductory programming classes, such as CIS 110 or 120 at the University of Pennsylvania. Features of interest included line length, modularity, and consistency. They used machine learning to compute a single score from the features. CodeScore will make use of similar techniques, but the two projects will be using mostly non-overlapping feature sets and different strategies for identifying them. Our focus is not on style but on code smells as they relate to internal quality, which ultimately affects the correctness, efficiency, and overall external quality of the program.

2.6 Other Static Code Analysis Tools

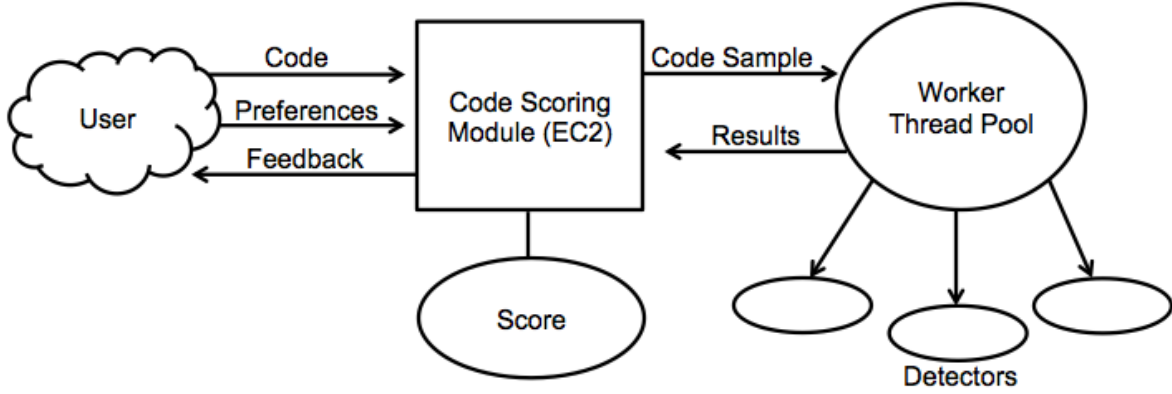A wide variety of other static code analysis tools also ex-

Figure 1: System architecture

ist. The Web Application Security Consortium provides a community-sourced list of some commonly used tools for code analysis [6]. Some of these tools include:

- CodeSonar - A full-featured code analysis tool with binary analysis, results visualization, and a wide variety of metrics. CodeSonar can detect numerous code faults and security risk within projects containing up to millions of lines of Java, C, or C++ code [9].

- Klocwork Insight - Provides a wide range of metrics related to security and reliability with detailed reports and offers integration with a variety of IDEs [12].

- FindBugs - An open-source code analysis tool which is capable of detecting many different types of bugs, such as null pointer dereferences or infinite loops, in Java code [21].

- Yasca - Another open-source code analysis tool combining the functionality of several similar tools in order to provide quality and security metrics for code written in a wide range of languages [22].

3. SYSTEM MODEL

CodeScore offers a novel solution amongst all of these existing tools. Most of the current applications are principally concerned with security, reliability, correct functionality, efficiency, and other aspects of external quality. Our tool focuses on internal quality, which can help prevent problems with external quality, thus allowing developers to spend less time testing and fixing bugs and more time on new products and features. Additionally, many of the existing tools target industry developers, but one of our primary target audiences will be students, who need to need to develop an appreciation for internal quality before entering the workforce or academia. Finally, our solution attempts to provide more functionality than the majority of open-source, similar applications. We aim to provide quality of functionality similar to commercial-grade applications, while making our solution accessible to a wider audience and not merely for large companies.

We have designed a program called CodeScore that quantitatively assesses code quality. It is an objective tool for performing a task that seems inherently subjective. However,

the characterization of software quality has been the topic of much discussion in both academia in industry. Using as a foundation the components of internal quality described in [11] and similar reports, our program further dissects these components into specific patterns that can be identified in source code. The holistic scoring system is built up from a large collection of these specific, detectable indicators to provide one comprehensive measure of quality.

CodeScore focuses on evaluating internal quality and makes use of static code analyses. This allows for assessment of source code without bias from different system architectures or environments. Code smells provide a powerful but simple indicator of internal quality because they are specific and detectable, they have a clear relationship with internal quality, and because there is existing work to use as a reference [17], [18]. The program currently focuses on just two important metrics of internal quality, understandability and maintainability.

Understandability is a measure of how easy a code sample is for a human to interpret. It can be estimated in part by the length of message chains (a code smell in which one method invokes another, which invokes another, and so on in a long one-line sequence of function calls), length of parameter lists, by determining what fraction of variable names are dictionary words versus strings of letters and numbers, and by analyzing the class structure of a program.

Maintainability is a measure of how easy a code sample will be to update and change. It can be estimated in part by detecting and recognizing coupling between classes, duplicated code, and classes that are too large. Another indicative code smell is shotgun surgery, which describes the situation in which changing one behavior requires numerous changes to code in different places. For example, if logging statements are implemented separately in each function in a class, then adding line numbers to the logs will require considerable time and effort. A better solution would be to write a log wrapper for all of the functions, so that any changes only need to be made once.

CodeScore implements the workflow illustrated in Figure 1. The key components of the system are the main driver and thread pool, the detectors, and the scoring module. First, a user uploads source code, an optional revision history from a version control repository, and a JSON file of

preferences to our server. Detectors for several code smells related to understandability and maintainability run in parallel on all files in the source code and the revision history. The results of the detections are then combined into an overall score for the code sample, analogous to a grade on an essay. The program also provides a detailed score report describing the number, type, and location of problems in the code. This report can be shared with recruiters as a quick and objective evaluation of a candidate's programming ability, used as a tool for self-improvement, or incorporated into teachers' grading rubrics.

The detectors are the most complicated component of the system and present the most technical challenges. Each code smell requires a unique detection algorithm, sometimes involving analysis of multiple classes together or complex parsing. The detectors currently implemented have been made successful by first creating a detailed list of all possible patterns to search for using the syntactical structure of the program, then implementing the search pattern and a counter to keep track of the number of violations. For any detectors that require some sort of threshold, like the maximum depth allowed for nested conditionals, the program's default parameters can be overriden using a preference configuration file.

Some code smells cannot easily be detected using the source code, but can be found by taking advantage of additional information. For instance, assume every time a method A.foo() is changed, several additional changes must be made elsewhere in the code. This issue would be evident upon careful inspection of the revision history but difficult to detect in the source code itself. This is one of the next features that we plan to incorporate into CodeScore. The challenging implementation is discussed in the next section. Adding revision history analyses to our list of metrics will help CodeScore offer a more well-rounded assessment of software projects, and they are high priorities among our next steps.

Another challenge is making the program efficient. CodeScore is designed to run on large and small software projects, which means that all parsing, processing, and reporting must be carefully engineered to produce results in a timely manner. CodeScore employs parallel processing techniques so that the runtime scales well with the number of metrics and the size of the input, and it takes advantage of efficient parsing APIs to assist with code smell detection.

4. SYSTEM IMPLEMENTATION

All processing for CodeScore happens on an Amazon EC2 cloud processing machine. EC2 is a service that provides resizable computing capacity in the cloud. Users upload Java code, a preferences file, and an optional revision history from a version control repository to the EC2 instance through a simple web app. The preferences file is encoded in JSON and includes information that is specific to each detector, such as thresholds for certain code smells. A sample preferences file can be seen in Figure 2. If no preferences file is includes, the program uses default values. The controller class and all detectors are implemented in Java. When software is uploaded to the endpoint, the controller class generates n worker threads, where n is the number of detection algorithms to be performed on the software sample. Currently, $n = 3$ since we support 3 detectors, but we hope to eventually implement 8 to 12 detectors. Each of

the workers performs a specific detection task on the uploaded files. Once complete, the controller then reduces all of the data computed by the workers and summarizes the findings. A maintainability score and an understandability score are computed using the results of the relevant code smell detection algorithms. The scores are informed by the frequency and severity of the issues (one poorly named variable is less problematic than five duplicated code blocks), and the sum of these two scores is reported as the overall CodeScore. The scoring algorithm has not been fully implemented, but it will assign a weight the results from each of the code smell detectors based on how strongly each code smell impacts internal code quality. As a starting point, these weights can be configured with the user preferences file. Later, these weights will be tuned using feedback from user studies, in which other people will be asked to rate the severity of various code smells with regard to their affect on internal software quality.

```json
{
    "preferences": {
        "longestMethodChain": 1,
        "maxConditionalDepth": 1,
        "maxParamCount": 3
    },
    "weights": {
        "methodChain": ".21",
        "nestedConditionals": ".49",
        "paramCounts": ".3"
    }
}
```

Figure 2: User preferences file encoded in JSON

The detectors use the Eclipse Java Abstract Syntax Tree (AST) API to parse the Java code as a first step in finding code smells. Traversing the AST allows the system to access and manipulate Java code at the syntactic level, efficiently searching for specific elements such as if statements or method declarations. Detecting code smells at the AST level will also make CodeScore easy to adapt to additional programming languages in the future. Most object-oriented code smells have the same properties from the perspective of the AST, which abstracts away the specifics of individual programming languages.

A sample abstract syntax tree and the Java code it represents can be seen in Figure 3.

We have prioritized the code smells and have implemented three of them so far. Some code smells are simpler to detect than others, and we currently have detectors for the following:

- Long parameter lists. Method declarations are parsed using the AST API, and then parameter lists are extracted to determine their length. The maximum length parameter list can be configured from the JSON preferences file.

- Deeply nested conditional logic. Conditional statements including if statements and loops are parsed using the AST API. Traversing the program as a tree results in finding parent-child relationships between nested conditional statements. The maximum nesting depth is configurable from the preferences file.
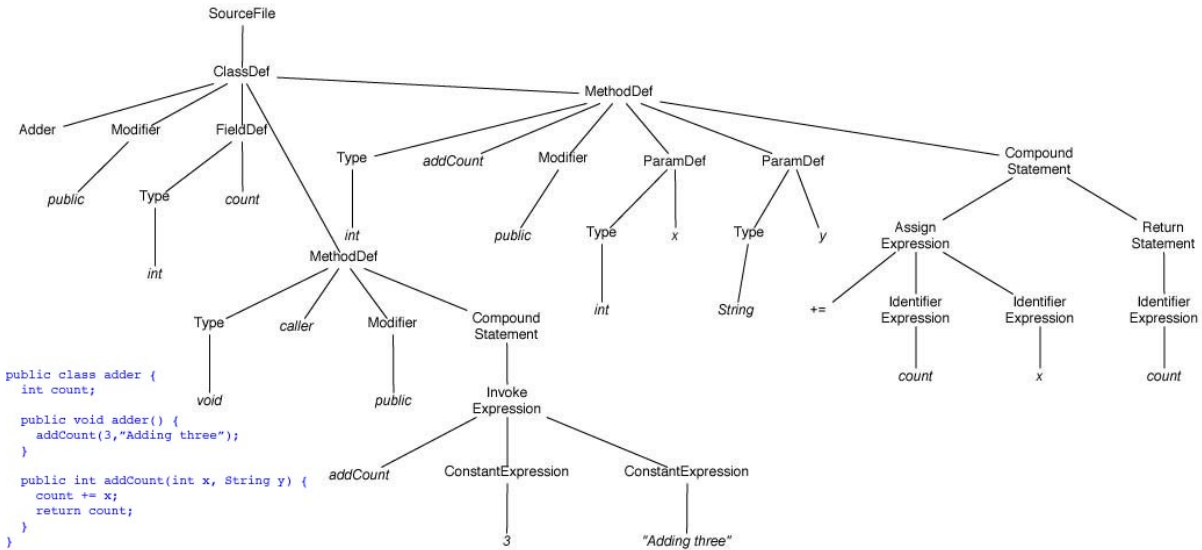
Figure 3: Sample Java abstract syntax tree

- Message chaining. Method invocations can be isolated using the AST API. By examining the context information provided by the AST API about each method invocation, it is possible to detect when multiple method calls are chained together. The maximum length of chaining to allow is configurable in the preferences file.

All of these code smells detract from understandability. Poor naming conventions also detract from understandability. We have implemented a basic natural language processing module using the Apache OpenNLP library. This library will be used to analyze variable and method names to determine if they are informative (containing some form of a dictionary word or words) and consistent. Variable and method declarations can be extracted from the source code using the AST API.

Code smells that affect both understandability and maintainability are long methods and large/complex classes. Both of these will be detected by finding the class or method declaration using the AST and analyzing the contents for length. Additionally, to measure class complexity, measures such as the McCabe Cyclomatic Complexity [**CITE THIS**] can be calculated.

Code smells that affect primarily maintainability are duplicated code, feature envy, and shotgun surgery. The Java Reflection API [**CITE THIS TOO**] faciliates duplicate code detection. Feature envy, indicative of an inappropriate relatinship between two classes, will require a more indepth analysis several structures in the AST. We can find references to each class field and method invocation within a particular class and determine if they correspond to a different class or the class being examined.

Shotgun surgery is a code smell affecting maintainability that requires more than just the source code to detect. As described previously, shotgun surgery is a problem in which changing one behavior requires several separate changes in the code. This antipattern can be detected by parsing revision histories to determine if certain methods or lines of code in different parts of the program tend to changed at the same time. Analyzing revision history will require the development of a custom parser that can extract and tokenize changes at a method-level granularity. We can do this by extracting a full change history (such as a git diff of two consecutive commits), removing the metadata symbols (for example, git uses plus symbols and minus symbols to indicate insertions and deletions) to extract just the code, and then using the Java AST API to examine the syntactic structure of the elements that were changed. The shotgun surgery detector will analyze the full history of the program, keeping track of which methods are modified together in order to see if a problematic pattern exists.

One drawback of the overall CodeScore system is that it is largely centralized. If the server containing the controller code is compromised at any time during the computation, all progress in the analysis will be lost. A reach goal or future development is to decentralize the process so that loss of a server does not result in complete loss of progress. Ideally, this could be implemented using MapReduce [7].

5. SYSTEM PERFORMANCE

At this stage in the implementation of CodeScore, there are two key aspects of our application that can be tested:

- How accurately code smells are detected.

- How quickly the analysis runs.

The detectors have been tested primarily with smaller code samples which were purposely polluted with the code smells we were seeking to detect: long parameter lists, deeply nested conditional logic, and message chaining. These code samples were analyzed by hand to construct an oracle from which we could judge the system's performance.

The system detected 100% of the code smells that were found by hand, which corresponds to 100% precision and recall. It should be noted that CodeScore is currently only equipped with three of the simpler code smell detectors, and we expect these scores to be lower, perhaps 80% after incorporating more complex detection algorithms.
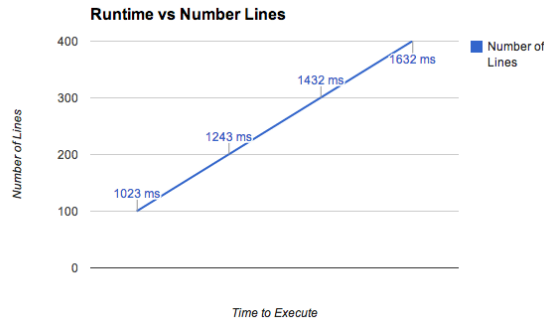
Figure 4: Preliminary Runtime Analysis of Code-Score Module

The speed of execution varied roughly linearly with the number of lines of code in preliminary tests. The bottleneck in runtime appears to be the Java Abstract Syntax Tree API. At this time, the benefits of using the AST API in terms of simplicity and modularity outweigh any added slowness in processing. However, we will investigate alternative methods if runtime with future detection algorithms becomes a serious issue. See Figure 4 for a graph of the runtime analysis results.

6. REMAINING WORK

We have a basic end-to-end system in place, so our focus now is to add features and expand our set of code smell detectors. The remaining steps are:

- Implement minimum of five more detection algorithms, including:
  - Duplicated code
  - Poor naming conventions
  - Feature envy
  - Shotgun surgery
  - Complex classes

- Create a scoring algorithm that computes a single score based on the findings of the detectors. Allow this scoring algorithm to be configured by a JSON file.
- Move the code so that it runs on EC2 instead of running locally. Create web front-end for users to upload their code.
- Perform user studies to compare CodeScore's grade to human evaluation. Incorporate feedback to improve detection and scoring algorithms.

Additionally, if the previous goals are accomplished earlier than expected, we will investigate how to change design to use map reduce. We will also look for more code quality metrics to incorporate and add additional code smells.

It is possible that once we have implemented all of the detectors, not all of the code smells will provide unique information. Some code smells might be so correlated with each other that performing the additional detection is not worth the increased processing time. If this is the case, we will discard one of the correlated smells and replace it with a more informative metric.

We estimate that we have completed 40% of the necessary tasks for CodeScore.

7. REFERENCES

[1] J Allen, S Collison, and R Luckey. Ohloh web site. `https://www.ohloh.net/`, 2009.

[2] Barry W Boehm, John R Brown, and Myron Lipow. Quantitative evaluation of software quality. In *Proceedings of the 2nd international conference on Software engineering*, pages 592–605. IEEE Computer Society Press, 1976.

[3] Barry W Boehm, Ray Madachy, Bert Steece, et al. *Software Cost Estimation with Cocomo II with Cdrom*. Prentice Hall PTR, 2000.

[4] Shyam R Chidamber and Chris F Kemerer. A metrics suite for object oriented design. *Software Engineering, IEEE Transactions on*, 20(6):476–493, 1994.

[5] Steve Christey. Plover: Preliminary list of vulnerability examples for researchers. In *NIST Workshop Defining the State of the Art of Software Security Tools*, 2005.

[6] The Web Application Security Consortium. Static code analysis tools. `http://projects.webappsec.org/w/page/61622133/StaticCodeAnalysisList`, 2012.

[7] Jeffrey Dean and Sanjay Ghemawat. Mapreduce: simplified data processing on large clusters. *Communications of the ACM*, 51(1):107–113, 2008.

[8] M Fowler. *Refactoring: Improving the design of existing code*. Addison-Wesley, 1999.

[9] GrammaTech. Codesonar. `http://www.grammatech.com/codesonar`, 2013.

[10] Maurice H Halstead. *Elements of Software Science*. Elsevier Science Inc., 1977.

[11] ISO ISO. Iec 25010: 2011: Systems and software engineering–systems and software quality requirements and evaluation (square)–system and software quality models. *International Organization for Standardization*, 2011.

[12] Klocwork. Klocwork insight. `http://www.klocwork.com/products/insight`, 2013.

[13] Nancy Leveson. *Safeware: System Safety and Computers*. Addison-Wesley, 1995.

[14] Thomas J McCabe. A complexity measure. *Software Engineering, IEEE Transactions on*, (4):308–320, 1976.

[15] Steve McConnell. *Code Complete*. Microsoft Press, 1993.

[16] CWE MITRE. Common weakness enumeration. `https://cwe.mitre.org/`, 2006.

[17] Naouel Moha, Yann-Gael Gueheneuc, Laurence Duchien, and Anne-Francoise Le Meur. Decor: A method for the specification and detection of code and design smells. `http://www.archipel.uqam.ca/5165/1/Moha09-TSE.pdf`, 2009.

[18] Fabio Palomba, Gabriele Bavota, Massimiliano Di Penta, Rocco Oliveto, Andrea De Lucia, and Denys Poshyvanyk. Detecting bad smells in source code using change history information. `http://www.cs.wm.edu/~denys/pubs/ASE%2713-HIST-Camera.pdf`.

[19] Robert E Park. Software size measurement: A framework for counting source statements. Technical report, DTIC Document, 1992.

[20] R Pressman. *Software Engineering: A Practicioner's Approach.* McGraw-Hill, 4 edition, 1997.

[21] Bill Pugh and Andrey Loskutov. Findbugs. http://findbugs.sourceforge.net, 2013.

[22] Michael Scovetta. Yasca. http://www.scovetta.com/yasca.html, 2007.